
What Hast Earth to do With Mars: An explanation of Urbit's IPC interfaces

~mopfel-winrux
Native Planet

Abstract

In this groundbreaking paper, we embark on a cosmic journey with the renowned character Ernest P. Worrell as he ventures into the unexplored realm of Martian computing. Drawing inspiration from Ernest’s comically ingenious encounters with everyday challenges, we investigate the foundations of what we term “Artificial Stupidity.” As Ernest grapples with Martian technology, we delve into the intricacies of programming errors, algorithmic missteps, and the curious phenomena that arise when human-like intelligence meets extraterrestrial computing systems. Our analysis sheds light on the unexpected intersections between humor, artificial intelligence, and the cosmic absurdity of Martian software. Join us in this interplanetary exploration as we unravel the mysteries of Artificial Stupidity through the lens of Ernest’s interstellar escapades.

Contents

1	Introduction	2
2	Urbit’s Unique Landscape	2
2.1	Arvo: The Heart of Martian Computing	2
2.2	Vere: The Substrate of Martian Technology	3
3	Khan - The High-Level Thread Interface of Urbit	4
4	conn.c	5
4.1	%ovum	5
4.2	%fyrd	6
4.3	%urth	6
4.4	%peek	6
4.5	%peel	7
5	%lick	7
5.1	/sys/lull Definition	8

32 1 Introduction

33 In the realm of software, two worlds exist in stark contrast: Earth, the familiar domain
 34 of established software paradigms, bustling with complex and varied digital ecosys-
 35 tems, and Mars, the enigmatic realm of Urbit, a system as pristine and archaic as the
 36 Martian landscape itself. This article sets forth on an expedition to unravel the mys-
 37 teries of this Martian software landscape, exploring how the intricate processes of
 38 terrestrial technology can establish a dialogue with the Martian code.

39 Unlike Earth, where software is a tapestry of evolution and patchwork, Mars
 40 presents a realm of purity. It is a world where software exists as a Maxwellian con-
 41 struct, untouched and unaltered by the non-Maxwellian intricacies that characterize
 42 Earth's digital environments. Yet, in this isolation lies a profound awareness - an un-
 43 derstanding by the Martians that there exist diverse forms of computing with which
 44 they must engage. Mars, in its wisdom, accommodates these external entities, allow-
 45 ing them to interact in a dialect familiar to its inhabitants: the language of nouns.

46 Our journey into this uncharted territory begins with a deep dive into the archi-
 47 tecture of Urbit, focusing on the foundational layer upon which Arvo, Urbit's kernel,
 48 operates. This exploration will illuminate the operational dynamics and philosophi-
 49 cal underpinnings that make Urbit a universe apart. Following this, we delve into
 50 the intricacies of two pivotal components of Urbit's ecosystem: %khan and %lick.
 51 These vanes serve as distinct conduits between the terrestrial and the Martian, facili-
 52 tating not just interaction but a nuanced form of control and collaboration with Arvo.
 53 Through this exploration, we aim to demystify the enigma of Urbit and bridge the
 54 cosmic gap between Earth's established software paradigms and the Martian vision
 55 of computational purity and simplicity.

56 2 Urbit's Unique Landscape

57 2.1 Arvo: The Heart of Martian Computing

58 Arvo, also known as Urbit OS, represents a paradigm shift in operating system design.
 59 Unlike traditional operating systems that operate on preemptive multitasking and
 60 complex event networks, Arvo is a purely functional operating system, characterized
 61 by its deterministic nature and compact size. The entire Urbit stack is about 30,000
 62 lines of code, with Arvo itself being only around 1,000 lines. This small codebase is
 63 intentional, reflecting a philosophy that system administration complexity is directly
 64 proportional to code size.

65 Arvo's design philosophy positions it not just as an operating system but as a new
 66 frontier in the peer-to-peer internet space. Unlike traditional operating systems like
 67 Windows, macOS, or Linux, Arvo doesn't aim to replace them; instead, it offers a
 68 unique user experience akin to a web browser, yet it is a fully-fledged OS in its own

69 right. This distinctive approach enables Arvo to function within a virtual machine,
70 theoretically capable of running on bare metal.

71 Arvo's unique architecture avoids what is referred to as "event spaghetti" by main-
72 taining a clear causal chain for every computation. Each chain begins with a Unix I/O
73 event and progresses through a series of steps until the computation's terminal cause.
74 This deterministic nature allows for a high level of predictability and control, a stark
75 contrast to the non-deterministic nature of most Earth-based operating systems.

76 The kernel's design as a "purely functional operating system" – or more accurately,
77 an "operating function" – underscores its uniqueness. Arvo operates on the principle
78 that the current state is a pure function of its event log, a record of every action ever
79 performed. This determinism is a significant deviation from the norm, where oper-
80 ating systems allow for programmatic alteration of global variables affecting other
81 programs.

82 Arvo handles nondeterminism in an innovative way. The system, in essence, be-
83 haves like a stateful packet transceiver – events are processed, but there's no guar-
84 antee of completion, mirroring the behavior of packet dropping in networking. This
85 approach to handling nondeterminism is unique and aligns with Arvo's overall phi-
86 losophy of simplicity and determinism.

87 Arvo's determinism is a key feature, stacking it atop a frozen instruction set known
88 as Nock. This concept, though new to operating systems, is not foreign to comput-
89 ing. Similar to how CPU instruction sets like x86-64 are frozen at the chip level, Arvo
90 freezes the instruction set at a higher level, enabling deterministic computation. This
91 high-level determinism contrasts sharply with the non-determinism typically found
92 in Earth's operating systems.

93 Handling nondeterminism in Arvo is akin to a heuristic decision-making process
94 similar to dropping a packet in networking. This approach views Arvo as a stateful
95 packet transceiver, where the completion of events is not guaranteed, a stark differ-
96 ence from traditional computing models. Additionally, because Arvo runs on a VM,
97 it can obtain nondeterministic information, such as stack traces from infinite loops,
98 through the interpreter beneath it, further enhancing its operational capabilities.

99 The vision of Urbit, with Arvo at its core, is to transition from developer-hosted
100 web services on multiple foreign servers to self-hosted applications on a personal
101 server. Each architectural decision in Arvo aligns with this vision, emphasizing user
102 ownership and management of data. This approach sets it apart from the conventional
103 model of cloud-based services and centralized data management prevalent in Earth-
104 based systems.

105 2.2 Vere: The Substrate of Martian Technology

106 Vere, the runtime of Urbit, plays a critical role in actualizing the Martian comput-
107 ing model on Earth-based hardware. As the Nock interpreter written in C, Vere is
108 intricately optimized to run Arvo, ensuring seamless translation of its deterministic
109 operations into practical execution on conventional systems.

110 Key to Vere's functionality is how it processes events. Events from the Unix en-
111 vironment are passed to Vere, which then translates and injects them into Arvo. This

mechanism demonstrates a profound integration between Vere and Arvo, where Vere possesses direct knowledge of Arvo's state. Notably, Vere can 'scry' or query Arvo's state without altering it, retrieving information as needed. This capability allows Vere to maintain the integrity of Arvo's deterministic model while enabling dynamic interaction with the outside world.

Urbit's use of Unix as its BIOS is a pivotal aspect of its runtime operation. By leveraging the robust, widely-used Unix system, Vere ensures that Urbit can run on a broad range of hardware platforms, effectively making Urbit's Martian technology accessible and operable in Earth's diverse computing environments. This strategic use of Unix not only provides the necessary I/O and optimizations for Arvo but also ensures that Urbit's advanced and unique software architecture can function in tandem with the existing, established hardware and software ecosystems of Earth.

Moreover, Vere's design allows it to inject events into Arvo, a feature crucial for maintaining the fluid communication and interaction between the Urbit system and its underlying hardware and software infrastructure. This bidirectional communication channel ensures that Arvo can respond to external stimuli while remaining true to its functional and deterministic nature.

Thus, Vere serves as more than just a bridge between Martian technology and Earthly hardware; it is a finely tuned conduit, optimized to uphold and facilitate the unique computational paradigm that Arvo embodies.

3 Khan - The High-Level Thread Interface of Urbit

Khan, as a crucial component of Urbit's architecture, functions as a high-level thread interface, pivotal for managing both internal and external communications within the Urbit ecosystem. It is designed to efficiently handle complex I/O operations, employing an IO monad coupled with an exception monad. This dual-mechanism approach allows Khan to adeptly manage complex IO tasks while effectively handling potential failures.

Khan's role in Urbit can be likened to that of a control plane. Its primary function is to execute threads via a Unix Socket and relay the results back. This setup supports various thread-running modes, enhancing the flexibility of Urbit's computing environment and allowing for diverse interaction methods with both internal and external processes.

At its core, Khan's design revolves around the use of 'threads', which are monadic functions designed to handle arguments and produce results. These threads are optimized for I/O operations, leveraging the IO monad for structured and efficient task handling. The exception monad further ensures robust management of potential failures and unexpected events.

Khan's implementation ensures that Urbit serves as a personal server, replacing multiple developer-hosted web services on foreign servers with self-hosted applications on a single personal server. Its design underscores Urbit's commitment to user ownership and data management, setting it apart from the conventional model of cloud-based services and centralized data management prevalent in Earth-based systems.

4 `conn.c`

`conn.c`¹ is a driver in Vere. It is a part of the "King" (a.k.a. "Urth") process. It exposes a Unix domain socket² at `/path/to/pier/.urb/conn.sock` for sending/receiving data from external processes.

The functionality of `conn.c`, particularly its ability to dispatch messages and handle various types of requests, is critical for the command of urbit's control plane. It serves as a way for unix processes to receive insights about what is happening on urbit. It also allows for pass commands either `%khan` and the injection of raw kernel moves into the system. The implication of this means that whoever is running your urbit runtime controls your urbit.

From a technical perspective, `conn.c` accepts newt-encoded ++ jammed nouns which take the form `[request-id command arguments]`. This structure allows for a variety of commands to be executed, including

1. `%ovum` - the injection of raw kernel moves
2. `%fyrd` - a direct shortcut to Khan commands
3. `%urth` - runtime subcommands like `%pack'` or `%meld`
4. `%peek` - namespace scry requests into Arvo
5. `%peel` - emulated namespace scry requests into Vere

A valid `conn.c` command produces a newt-encoded jammed noun with type `[request-id output]`, where:

- `request-id` matches the input `request-id`
- `output` depends on the command

An invalid `conn.c` command produces a newt-encoded jammed noun with type `[0 %bail error-code error-string]`.

4.1 `%ovum`

The argument to an `%ovum` command is a raw kernel move which is injected directly into the Arvo event loop. This is a very powerful - and potentially dangerous - tool. For example, if a ship somehow got into a state where Clay was no longer working properly (meaning new files could not be compiled to fix the state of the kernel), the source code for a new, working Clay could be directly injected into the ship using an `%ovum`.

The output of an `%ovum` command is:

- `[%news %done]` if the move completed successfully
- `[%news %drop]` if the move was dropped
- `[%bail goof]` if an error occurred

¹<https://github.com/urbit/vere/blob/develop/pkg/vere/io/conn.c>

²https://en.wikipedia.org/wiki/Unix_domain_socket

4.2 %fyrd

%fyrd is a direct shortcut to the Khan vane. The arguments to a %fyrd command are (in order):

1. The name of the desk in which the thread lives (e.g. %base) or beak for the thread (e.g. [%zod %base %10])
2. The name of the thread (e.g. %hi)
3. Mark to which the output should be cast (e.g. %tape)
4. Mark for how to interpret the input argument to thread (e.g. %ship)
5. Input argument to thread (e.g. zōd)

The output of a %fyrd command is [%avow (each page goof)], the value of each depending on whether the thread succeeded or not.

4.3 %urth

The argument to the %urth command is a subcommand for the action to perform. Currently, the only valid commands are %pack and %meld.

%urth will return %& if given a valid command as input, otherwise it will return [0 %bail 0xffffffff9 %urth-bad]. No other output is emitted.

4.4 %peek

The %peek command is used to perform a namespace read request (a.k.a. scry) using Arvo's external peek interface arm +22 in arvo.hoon³. The argument to %peek is the nom input to +peek in arvo.hoon (lyc is auto-filled as [], i.e., "request from self"). That is to say that the argument to %peek must have type:

```
$+ each path
$% [%once vis=view syd=desk tyl=spur]
   [%beam vis=view bem=beam]
==
```

Practically speaking, this means that the input will look like one of these three examples:

```
[%& p=path]
[%| p=[%once vis=view syd=desk tyl=spur]]
[%| p=[%beam vis=view bem=beam]]
```

Where:

- path is a [view beam], with the view passed in as a coin

³<https://github.com/urbit/urbit/blob/develop/pkg/arvo/sys/arvo.hoon#L1774>

- 222 • `view` is the vane code for the scry, as well as an optional care, possibly appended
223 to the vane (e.g. `%j`, `%gx`, etc.)
- 224 • `beam` is a `[beak spur]`
- 225 • `desk` is used to auto-generate a `beak`: `[our desk now]`
- 226 • `spur` is the scry endpoint for the agent or vane

227 The output of a `%peek` command is `[%peek (unit (unit scry-output))]`,
228 where means that the scry endpoint is invalid, and `[]` means that the scry
229 resolved to nothing.

230 4.5 %peel

231 `%peel` attempts to emulate a scry-like namespace, like the one used by Arvo and
232 accessed by `%peek`. The argument to `%peel` should be a path. Valid paths result in
233 a non-null `unit` containing the result of the scry. Invalid paths result in null (i.e.).
234 The valid paths and the data they return are:

235	<code>/help</code>	<code>(unit (list path))</code>	Supported <code>%peel</code> paths
236	<code>/live</code>	<code>(unit ~)</code>	Pier health check; succeeds if pier is running
237	<code>/khan</code>	<code>(unit ~)</code>	Khan health check; succeeds if Khan vane is running
238	<code>/info</code>	<code>(unit mass)</code>	Pier info as a mass
239	<code>/v</code>	<code>(unit @t)</code>	Returns version of the Vere binary as a cord
240	<code>/who</code>	<code>(unit @)</code>	Returns the Azimuth identity of the ship as an atom

241 Note that the pier info above is returned as a `mass` report, i.e. type `(pair cord`
242 `(each * (list mass)))`. This is not the same as the `|mass` memory report.
243 `/mass` is meant to be a valid `%peel` path which returns the `|mass` memory report,
244 but it is currently unimplemented.

245 5 %lick

246 Altho also dealing with interprocess communication, Lick was designed for a very dif-
247 ferent scenario than Khan: to allow external processes, in particular hardware drivers,
248 to intercommunicate with Urbit. (This breached the Earth/Mars divide.) Thus `/sys/-`
249 `vane/lick` focuses on instrumenting a low-level noun interfaces over domain sock-
250 ets.

251 `%lick` manages IPC ports, and the communication between Urbit applications
252 and POSIX applications via these ports. Other vanes and applications ask `%lick` to
253 open an IPC port, notify it when something is connected or disconnected, and transfer
254 data between itself and the Unix application.

255 Lick works by opening a Unix socket for a particular process, which allows seri-
256 alized IPC communications. These involve a jammed noun so the receiving process
257 needs to know how to communicate in nouns.

258 The IPC ports Lick creates are Unix domain sockets (AF_UNIX address family) of
 259 the SOCK_STREAM type.

260 The connexions are made via filepaths in .urb/dev of the pier.

261 The format is:

262 V.BBBB.JJJJ.JJJJ...

- 263 • V is the version
- 264 • B is the jam size in bytes (little endian)
- 265 • J is the jammed noun (little endian)

266 The process on the host OS must therefore strip the first 5 bytes, ++cue⁴ the
 267 jamfile, check the mark and (most likely) convert the noun into a native data structure.

268 5.1 /sys/lull Definition

```

269 ::                                     ::::
270 ::::                                ++lick      :: (1j) IPC
271 ::                                     ::::
272 ++ lick ^?
273 |%
274 +$ gift                                :: out result <-$
275   $% [%spin =name]                    :: open an IPC port
276     [%shut =name]                     :: close an IPC port
277     [%spit =name =mark =noun]         :: spit a noun to the IPC port
278     [%soak =name =mark =noun]         :: soak a noun from the IPC port
279 ==
280 +$ task                                :: in request ->$
281   $~ [%vega ~]                         ::
282   $% $>(%born vane-task)               :: new unix process
283     $>(%trim vane-task)                 :: trim state
284     $>(%vega vane-task)                 :: report upgrade
285     [%spin =name]                       :: open an IPC port
286     [%shut =name]                       :: close an IPC port
287     [%spit =name =mark =noun]           :: spit a noun to the IPC port
288     [%soak =name =mark =noun]           :: soak a noun from the IPC port
289 ==
290 ::
291 +$ name path
292 -- ::lick
  
```

293 To evaluate what /sys/vane/lick is doing, we need to look at Unix's IPC model
 294 briefly. IPC ("interprocess communication") describes any way that two processes in

⁴<https://docs.urbit.org/language/hoon/reference/stdlib/2p#cue>

an operating system’s shared context have to communicate with each other. `%lick` focuses on Unix domain sockets⁵, which are just communication endpoints⁶.

For instance, a valid use of `%lick` would use cards that look like this:

```

++ init [[%pass / %arvo %l %spin /control]~ this]
::
++ on-arvo
|= [=wire =sign-arvo]
?+ sign-arvo (on-arvo: def wire sign-arvo)
  [%lick %soak *]
  ?+ mark.sign-arvo [~ this]
  ::
  %connect
  ~& > "connect"
  :_ this [%pass /spit %arvo %l %spit /control %init area.state]~
== ==
::
++ send-state
|= =state
^- card:agent:gall
[%pass /spit %arvo %l %spit /control %state [slick:state face.state food.state live.state]]

```

The vane definition of `/sys/vane/lick` is even simpler than `/sys/vane/khan`: it has no `++abet` core and primarily communicates to the `unix-duct` in its state. The owner is a duct to handle the return `%soak`.

Gall needs to wrap `%soak` and `%spit` to route properly. See e.g. `%%+ap-generic-take`. This lets multiple agents share sockets with the same name, and each agent can have its own folder.

6 Conclusion

To summarize, why did you write it? Why do we care? What impact should it have on Urbit development?

Your bibliography is a separate BibTeX file. We use the `plainnat` bibliography style. You can use `natbib` citation commands like `\citep{wikipedia}` for parenthesized references. Use `\citet{wikipedia}` for inline references. You can also use `\citeauthor{wikipedia}` for the principal author’s name.

“You can use traditional TeX—or LaTeX—representations” [Varney, 1987].

“Or you can use fancy quotes—and symbols.”

⁵https://en.wikipedia.org/wiki/Unix_domain_socket

⁶<https://man7.org/linux/man-pages/man7/unix.7.html>

References

- 330
- 331 Jim Varney. Ernest goes to camp. Film, 1987. URL [https://www.imdb.com/title/](https://www.imdb.com/title/tt0092974/)
332 [tt0092974/](https://www.imdb.com/title/tt0092974/). Directed by John R. Cherry III. Produced by Stacy Williams. Written
333 by John R. Cherry III and Coke Sams. Starring Jim Varney, Victoria Racimo, John
334 Vernon, Iron Eyes Cody, Lyle Alzado, Gailard Sartain, Daniel Butler, Patrick Day,
335 Scott Menville, Jacob Vargas, and Todd Loyd. Buena Vista Pictures.