

---

# What Hast Earth to do With Mars: An explanation of Urbit's IPC interfaces

~mopfel-winx  
Native Planet

## Abstract

This paper explores the innovative approaches undertaken by developers to enable effective communication between the Urbit system and conventional Earth-based software systems. Initially, Urbit's integration with external systems relied on the use of two specific vanes, %eyre and %iris, serving as HTTP client and server respectively. This setup, however, presented challenges to the ideal of autonomous operation envisaged for Urbit, traditionally characterized by the principle that "Earth calls Mars; Mars doesn't call Earth."

Addressing these challenges, the paper introduces the development of two new vanes, %khan and %lick, which embody the core principle of uncontaminated communication. These advancements allow for a refined interaction model where Earth-based systems communicate with Urbit in a Urbit's own language, nouns.

The implications of these developments are significant, offering a novel framework for the integration of distinct computational ecosystems. This paper discusses the technical and conceptual underpinnings of these vanes, their development process, and the broader implications for Urbit's interaction with conventional software systems.

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Urbit's Unique Landscape</b>	<b>2</b>
2.1	Arvo: The Heart of Martian Computing . . . . .	2
2.2	Vere: The Substrate of Martian Technology . . . . .	3
<b>3</b>	<b>%khan - The High-Level Thread Interface of Urbit</b>	<b>4</b>
<b>4</b>	<b>conn.c</b>	<b>5</b>
<b>5</b>	<b>%lick - The Low Level IPC Interface</b>	<b>6</b>

Manuscript submitted for review.

Address author correspondence to ~mopfel-winx.

## 31 1 Introduction

32 In the realm of software, two worlds exist in stark contrast: Earth, the familiar domain  
33 of established software paradigms, bustling with complex and varied digital ecosys-  
34 tems, and Mars, the enigmatic realm of Urbit, a system as pristine and archaic as the  
35 Martian landscape itself. This article sets forth on an expedition to unravel the mys-  
36 teries of this Martian software landscape, exploring how the intricate processes of  
37 terrestrial technology can establish a dialogue with the Martian code.

38 Unlike Earth, where software is a tapestry of evolution and patchwork, Mars  
39 presents a realm of purity. It is a world where software exists as a Maxwellian con-  
40 struct, untouched and unaltered by the non-Maxwellian intricacies that characterize  
41 Earth’s digital environments. Yet, in this isolation lies a profound awareness - an un-  
42 derstanding by the Martians that there exist diverse forms of computing with which  
43 they must engage. Mars, in its wisdom, accommodates these external entities, allow-  
44 ing them to interact in a dialect familiar to its inhabitants: the language of nouns.

45 In the original explanation of of Urbit, [Yarvin2010](#), explicitly says “The general  
46 structure of cross-planet computation is that Earth always calls Mars; Mars never  
47 calls Earth. The latter would be quite impossible, since Earth code is non-Maxwellian.  
48 There is only one way for a Maxwellian computer to run non-Maxwellian code: in a  
49 Maxwellian emulator.” This however has never been true for Urbit as it was imple-  
50 mented. Urbit has two vanes `%eyre` and `%iris` that are written as a specific interface  
51 to Earth (a web server and curl, respectively). One can even make the argument that  
52 Urbit running on a virtual machine specifically breaks this convention.

53 Our journey into this uncharted territory begins with a deep dive into the archi-  
54 tecture of Urbit, focusing on the foundational layer upon which Arvo, Urbit’s kernel,  
55 operates. This exploration will illuminate the operational dynamics and philosophi-  
56 cal underpinnings that make Urbit a universe apart. Following this, we delve into  
57 the intricacies of two pivotal components of Urbit’s ecosystem: `%khan` and `%lick`.  
58 These vanes serve as distinct conduits between the terrestrial and the Martian, facili-  
59 tating not just interaction but a nuanced form of control and collaboration with Arvo.  
60 Through this exploration, we aim to demystify the enigma of Urbit and bridge the  
61 cosmic gap between Earth’s established software paradigms and the Martian vision  
62 of computational purity and simplicity.

## 63 2 Urbit's Unique Landscape

### 64 2.1 Arvo: The Heart of Martian Computing

65 Arvo, also known as Urbit OS, represents a paradigm shift in operating system design.  
66 Unlike traditional operating systems that operate on preemptive multitasking and  
67 complex event networks, Arvo is a purely functional operating system, characterized  
68 by its deterministic nature and compact size. The entire Urbit stack is about 80,000

69 lines of code, with Arvo itself being only around 2,000 lines. This small codebase is  
 70 intentional, reflecting a philosophy that system administration complexity is directly  
 71 proportional to code size.

72 Arvo's unique architecture avoids what is referred to as "event spaghetti" by main-  
 73 taining a clear causal chain for every computation. Each chain begins with a Unix I/O  
 74 event and progresses through a series of steps until the computation's terminal cause.  
 75 This deterministic nature allows for a high level of predictability and control, a stark  
 76 contrast to the non-deterministic nature of most Earth-based operating systems.

77 The kernel's design as a "purely functional operating system" – or more accurately,  
 78 an "operating function" – underscores its uniqueness. Arvo operates on the principle  
 79 that the current state is a pure function of its event log, a record of every action ever  
 80 performed. This determinism is a significant deviation from the norm, where oper-  
 81 ating systems allow for programmatic alteration of global variables affecting other  
 82 programs.

83 Arvo handles nondeterminism in an innovative way. The system, in essence, be-  
 84 haves like a stateful packet transceiver – events are processed, but there's no guar-  
 85 antee of completion, mirroring the behavior of packet dropping in networking. This  
 86 approach to handling nondeterminism is unique and aligns with Arvo's overall phi-  
 87 losophy of simplicity and determinism.

88 Arvo's determinism is a key feature, stacking it atop a frozen instruction set known  
 89 as Nock. This concept, though new to operating systems, is not foreign to comput-  
 90 ing. Similar to how CPU instruction sets like x86-64 are frozen at the chip level, Arvo  
 91 freezes the instruction set at a higher level, enabling deterministic computation. This  
 92 high-level determinism contrasts sharply with the non-determinism typically found  
 93 in Earth's operating systems.

94 Handling nondeterminism in Arvo is akin to a heuristic decision-making process  
 95 similar to dropping a packet in networking. This approach views Arvo as a stateful  
 96 packet transceiver, where the completion of events is not guaranteed, a stark differ-  
 97 ence from traditional computing models. Additionally, because Arvo runs on a VM,  
 98 it can obtain nondeterministic information, such as stack traces from infinite loops,  
 99 through the interpreter beneath it, further enhancing its operational capabilities.

100 The vision of Urbit, with Arvo at its core, is to transition from developer-hosted  
 101 web services on multiple foreign servers to self-hosted applications on a personal  
 102 server. Each architectural decision in Arvo aligns with this vision, emphasizing user  
 103 ownership and management of data. This approach sets it apart from the conventional  
 104 model of cloud-based services and centralized data management prevalent in Earth-  
 105 based systems.

## 106 2.2 Vere: The Substrate of Martian Technology

107 Vere, the runtime of Urbit, plays a critical role in actualizing the Martian comput-  
 108 ing model on Earth-based hardware. As the Nock interpreter written in C, Vere is  
 109 intricately optimized to run Arvo, ensuring seamless translation of its deterministic  
 110 operations into practical execution on conventional systems.

111 The architecture of Vere is bifurcated into two distinct parts: the `king` and the `serf`.

The `king` component is responsible for handling vane I/O and managing the effects on Earth. In contrast, the `serf` operates as the process running the Arvo virtual machine (VM). This division of responsibilities allows for a clear delineation of tasks within the Vere runtime, enhancing its efficiency and robustness.

Communication between the `king` and `serf` is achieved through an Inter-Process Communication (IPC) port. This IPC mechanism facilitates a continuous dialogue between the two components, ensuring that the `king` can effectively manage external interactions while the `serf` maintains the integrity of the Arvo VM.

A core opcode in Urbit's assembly language (`nock`) is the `hint` opcode (opcode 11) this allows Urbit to pass hints to Vere which produces effects on Earth. This bridges the Earth/Mars gap and allows for vanes to interact outside of Mars.

Key to Vere's functionality is how it processes events. Events from the Unix environment are passed to Vere, which then translates and injects them into Arvo. This mechanism demonstrates a profound integration between Vere and Arvo, where Vere possesses direct knowledge of Arvo's state. Notably, Vere can 'scry' or query Arvo's state without altering it, retrieving information as needed. This capability allows Vere to maintain the integrity of Arvo's deterministic model while enabling dynamic interaction with the outside world.

Urbit's use of Unix as its BIOS is a pivotal aspect of its runtime operation. By leveraging the robust, widely-used Unix system, Vere ensures that Urbit can run on a broad range of hardware platforms, effectively making Urbit's Martian technology accessible and operable in Earth's diverse computing environments. This strategic use of Unix not only provides the necessary I/O and optimizations for Arvo but also ensures that Urbit's advanced and unique software architecture can function in tandem with the existing, established hardware and software ecosystems of Earth.

Moreover, Vere's design allows it to inject events into Arvo, a feature crucial for maintaining the fluid communication and interaction between the Urbit system and its underlying hardware and software infrastructure. This bidirectional communication channel ensures that Arvo can respond to external stimuli while remaining true to its functional and deterministic nature.

Thus, Vere serves as more than just a bridge between Martian technology and Earthly hardware; it is a finely tuned conduit, optimized to uphold and facilitate the unique computational paradigm that Arvo embodies.

### 3 %khan - The High-Level Thread Interface of Urbit

Khan, as a crucial component of Urbit's architecture, functions as a high-level thread interface, pivotal for managing both internal and external communications within the Urbit ecosystem. It is designed to efficiently handle complex I/O operations, employing an IO monad coupled with an exception monad. This dual-mechanism approach allows Khan to adeptly manage complex IO tasks while effectively handling potential failures.

Khan allows threads to be triggered from outside of Urbit. A thread is a monadic function that takes arguments and produces a result. It may perform input and output while running, so it is not a pure function. A thread's strength is that it can easily

perform complex IO operations. It uses what's often called the IO monad (plus the exception monad) to provide a natural framework for IO. A thread's weakness is that it's impermanent and may fail unexpectedly. In most of its intermediate states, it expects only a small number of events (usually one), so if it receives anything it didn't expect, it fails. When code is upgraded, it's impossible to upgrade a running thread, so it fails.

Khan's role in Urbit can be likened to that of a control plane. Its primary function is to execute threads via a Unix Socket and relay the results back. This setup supports various thread-running modes, enhancing the flexibility of Urbit's computing environment and allowing for diverse interaction methods with both internal and external processes.

At its core, Khan's design revolves around the use of 'threads', which are monadic functions designed to handle arguments and produce results. These threads are optimized for I/O operations, leveraging the IO monad for structured and efficient task handling. The exception monad further ensures robust management of potential failures and unexpected events.

Khan's implementation ensures that Urbit serves as a personal server, replacing multiple developer-hosted web services on foreign servers with self-hosted applications on a single personal server. Its design underscores Urbit's commitment to user ownership and data management, setting it apart from the conventional model of cloud-based services and centralized data management prevalent in Earth-based systems.

## 4 conn.c

`conn.c`<sup>1</sup> is a driver in Vere. It is a part of the "King" (a.k.a. "Urth") process. It exposes a Unix domain socket<sup>2</sup> at `/path/to/pier/.urb/conn.sock` for sending/receiving data from external processes.

The functionality of `conn.c`, particularly its ability to dispatch messages and handle various types of requests, is critical for the command of urbit's control plane. It serves as a way for unix processes to receive insights about what is happening on urbit. It also allows for pass commands either `%khan` and the injection of raw kernel moves into the system. The implication of this means that whoever is running your urbit runtime controls your urbit.

From a technical perspective, `conn.c` accepts newt-encoded ++jammed nouns which take the form `[request-id command arguments]`. The newt-encoded format is:

`V.BBBB.JJJJ.JJJJ...`

- V is the version
- B is the jam size in bytes (little endian)
- J is the jammed noun (little endian)

<sup>1</sup><https://github.com/urbit/vere/blob/develop/pkg/vere/io/conn.c>

<sup>2</sup>[https://en.wikipedia.org/wiki/Unix\\_domain\\_socket](https://en.wikipedia.org/wiki/Unix_domain_socket)

193 This structure allows for a variety of commands to be executed, including

- 194 1. `%ovum` - the injection of raw kernel moves
- 195 2. `%fyrd` - a direct shortcut to Khan commands
- 196 3. `%urth` - runtime subcommands like `%pack` or `%meld`
- 197 4. `%peek` - namespace scry requests into Arvo
- 198 5. `%peel` - emulated namespace scry requests into Vere

199 A valid `conn.c` command produces a newt-encoded jammed noun with type `[request-`  
 200 `id output]`, where:

- 201 • `request-id` matches the input `request-id`
- 202 • `output` depends on the command

203 An invalid `conn.c` command produces a newt-encoded jammed noun with type  
 204 `[0 %bail error-code error-string]`.

## 205 5 %lick - The Low Level IPC Interface

206 Although also dealing with interprocess communication, `%lick` (**UIP-101**) was de-  
 207 signed for a very different scenario than Khan: to allow external processes, in partic-  
 208 ular hardware drivers, to intercommunicate with Urbit. (This breached the Earth/Mars  
 209 divide.) Thus `/sys/vane/lick` focuses on instrumenting a low-level noun interfaces  
 210 over domain sockets.

211 `%lick` manages IPC ports, and the communication between Urbit applications and  
 212 POSIX applications via these ports. Other vanes and applications ask `%lick` to open  
 213 an IPC port, notify it when something is connected or disconnected, and transfer data  
 214 between itself and the Unix application. Lick works by opening a Unix socket for  
 215 a particular process, which allows serialized IPC communications. These involve a  
 216 jammed noun so the receiving process needs to know how to communicate in nouns.  
 217 The IPC ports Lick creates are Unix domain sockets (AF\_UNIX address family) of the  
 218 SOCK\_STREAM type. The connexions are made via filepaths in `.urb/dev` of the pier.

219 The process on the host OS must therefore strip the first 5 bytes, `++cue`<sup>3</sup> the jamfile,  
 220 check the mark and (most likely) convert the noun into a native data structure.

221 To understand what `%lick` is doing, we need to look at Unix's IPC model briefly.  
 222 IPC ("interprocess communication") describes any way that two processes in an op-  
 223 erating system's shared context have to communicate with each other. `%lick` focuses  
 224 on Unix domain sockets<sup>4</sup>, which are just communication endpoints<sup>5</sup>.

225 For instance, a valid use of `%lick` would use cards that look like this:

<sup>3</sup><https://docs.urbit.org/language/hoon/reference/stdlib/2p#cue>

<sup>4</sup>[https://en.wikipedia.org/wiki/Unix\\_domain\\_socket](https://en.wikipedia.org/wiki/Unix_domain_socket)

<sup>5</sup><https://man7.org/linux/man-pages/man7/unix.7.html>

```

226 ++ init [[%pass / %arvo %l %spin /control]~ this]
227 ::
228 ++ on-arvo
229   |= [=wire =sign-arvo]
230   ?+ sign-arvo (on-arvo:def wire sign-arvo)
231     [%lick %soak *]
232     ?+ mark.sign-arvo [~ this]
233     ::
234     %connect
235     ~& > "connect"
236     :- this
237     [%pass /spit %arvo %l %spit /control %init area.state]~
238     ==
239   ==
240   ::
241   ++ send-state
242   |= =state
243   ^- card:agent:gall
244   :* %pass
245     /spit
246     %arvo %l
247     %spit
248     /control
249     %state
250     [slick:state face.state food.state live.state]
251   ==

```

252 The vane definition of `/sys/vane/lick` is even simpler than `/sys/vane/khan`: it  
 253 has no `++abet` core and primarily communicates to the `unix-duct` in its state. The  
 254 `+$owner` is a `+$duct` to handle the return `%soak`.

255 Arvo will send three types of `%soaks` to an open `%lick` port. These `%soaks` are  
 256 `%connect` when the first connection is established on the Earth-side IPC port, `%dis-`  
 257 `connect` when the last connection is broken on the Earth-side IPC port, and `%error`  
 258 when an error occurs. `%lick` also will send a `%disconnect` `%soak` to every agent when  
 259 Vere is started.

260 Gall needs to wrap `%soak` and `%spit` to route properly. See e.g. `%%+ap-generic-`  
 261 `take`. This lets multiple agents share sockets with the same name, and each agent can  
 262 have its own folder.

## 263 6 Conclusion

264 As Urbit had to deal with the reality of running on Earth software, Developers had  
 265 to come up with various ways of communicating between these systems. Originally  
 266 the two vanes `%eyre` and `%iris` were used as an HTTP client and server, breaking the  
 267 “Earth calls Mars; Mars doesn’t call Earth” dynamic. Recently, two other vanes, `%khan`

268 and %lick have been developed that hold true to the necessity idea of communication  
269 between Earth and Mars. These vanes force Earth to speak to Mars in a language it  
270 can understand, nouns.