
1 **What Hast Earth to do With Mars: An explanation of**
2 **Urbit's IPC interfaces**

3 **~mopfel-winxux**
4 **Native Planet**

5 **Abstract**

6 In the evolving landscape of digital computation, the interaction between con-
7 ventional software paradigms and innovative systems like Urbit presents a unique
8 challenge. This paper delves into the architectural intricacies and interfaces of Ur-
9 bit, contrasting its pristine, Maxwellian principles against the complex, evolved
10 nature of Earth's software ecosystems. Through a comprehensive examination of
11 Urbit's kernel (Arvo) and runtime (Vere), alongside its high-level and low-level
12 IPC interfaces (%khan and %lick, respectively), we elucidate the mechanisms that
13 facilitate communication and control between these disparate computing realms.
14 Our analysis aims to demystify Urbit's architecture and underscore its potential
15 to bridge the gap between the familiar digital environments of Earth and the un-
16 tapped possibilities of Martian computing.

17 **Contents**

18	1 Introduction	2
19	2 Urbit's Unique Landscape	2
20	2.1 Arvo: The Heart of Martian Computing	2
21	2.2 Vere: The Substrate of Martian Technology	3
22	3 %khan - The High-Level Thread Interface of Urbit	4
23	4 conn.c	5
24	5 %lick - The Low Level IPC Interface	6
25	6 Conclusion	7

1 Introduction

In the realm of software, two worlds exist in stark contrast: Earth, the familiar domain of established software paradigms, bustling with complex and varied digital ecosystems, and Mars, the enigmatic realm of Urbit, a system as pristine and archaic as the Martian landscape itself. This article will explore how Earth software establish a communication with the Martian code.

Unlike Earth, where software is a tapestry of evolution and patchwork, Mars presents a realm of purity. It is a world where software exists as a Maxwellian construct, it is defined by a small set of axioms (**Nock4k**). Urbit code is untouched and unaltered by the non-Maxwellian intricacies that characterize Earth's digital environments. Yet, in this isolation lies a profound awareness - an understanding by the Martians that there exist diverse forms of computing with which they must engage. Mars, accommodates these external entities by allowing them to interact in the Martian language of nouns.

In the original explanation of of Urbit, **Yarvin2010**, explicitly says "The general structure of cross-planet computation is that Earth always calls Mars; Mars never calls Earth. The latter would be quite impossible, since Earth code is non-Maxwellian. There is only one way for a Maxwellian computer to run non-Maxwellian code: in a Maxwellian emulator." This however has never been true for Urbit as it was implemented. Urbit has two vanes (best thought of as kernel modules) `%eyre` and `%iris` that are written as a specific interface to Earth (a web server and client, respectively). One can even make the argument that Urbit running on a virtual machine specifically breaks this convention.

This paper begins with a brief explanation of Urbit's architecture, focusing on Arvo (Urbit's kernel) and Vere (Urbit's Runtime). Following this, is an explanation of three core parts of the Urbit OS: `%khan`, `conn.c` and `%lick`. These serve as distinct conduits between the terrestrial and the Martian, facilitating not just interaction but a nuanced form of command and control of the entire system. Through this exploration, we aim to demystify the architecture of Urbit and bridge the gap between Earth's established software paradigms and the Martian vision of computation.

2 Urbit's Unique Landscape

2.1 Arvo: The Heart of Martian Computing

Arvo, also known as Urbit OS, represents a paradigm shift in operating system design. Unlike traditional operating systems that operate on preemptive multitasking and complex event networks, Arvo is a purely functional operating system, characterized by its deterministic nature and compact size. The entire Urbit stack is about 80,000 lines of code, with Arvo itself being only around 2,000 lines. This small codebase is intentional, reflecting a philosophy that system administration complexity is directly proportional to code size.

Arvo's unique architecture avoids what is referred to as "event spaghetti" by maintaining a clear causal chain for every computation. Each chain begins with a Unix I/O

67 event and progresses through a series of steps until the computation either completes
68 successfully or fails. This deterministic nature allows for a high level of predictability
69 and control, a stark contrast to most Earth-based operating systems.

70 The kernel’s design as a “purely functional operating system” – or more accu-
71 rately, an “operating function” – underscores its uniqueness. Arvo operates on the
72 principle that the current state is a pure function of its event log, a record of every
73 action ever performed. This determinism is a significant deviation from the norm,
74 where operating systems allow for programmatic alteration of global variables affect-
75 ing other programs.

76 Arvo handles non-determinism in an innovative way. The system, in essence, be-
77 haves like a stateful packet transceiver – events are processed, but there’s no guaran-
78 tee of the event successfully completing. This mirrors the behavior of packet dropping
79 in networking. This approach to handling non-determinism is unique and aligns with
80 Arvo’s overall philosophy of simplicity and determinism. Additionally, because Arvo
81 runs on a VM (Vere), it can obtain non-deterministic information, such as stack traces
82 from infinite loops, through the interpreter beneath it injecting events.

83 2.2 Vere: The Substrate of Martian Technology

84 Vere, the runtime of Urbit, plays a critical role in actualizing the Martian comput-
85 ing model on Earth-based hardware. As the Nock interpreter written in C, Vere is
86 intricately optimized to run Arvo, ensuring seamless translation of its deterministic
87 operations into practical execution on conventional systems.

88 The architecture of Vere is split into two distinct parts: the `king` and the `serf`.
89 The `king` component is responsible for handling vane I/O and managing the effects on
90 Earth. In contrast, the `serf` operates as the process running the Arvo virtual machine
91 (VM). This division of responsibilities allows for a clear delineation of tasks within the
92 Vere runtime, enhancing its efficiency and robustness.

93 Communication between the `king` and `serf` is achieved through an Inter-Process
94 Communication (IPC) port. This IPC mechanism facilitates a continuous dialogue
95 between the two components, ensuring that the `king` can effectively manage external
96 interactions while the `serf` maintains the integrity of the Arvo VM.

97 A core opcode in Urbit’s assembly language (Nock) is the `hint` opcode (opcode 11).
98 The `hint` opcode passes some value to the interpreter. The interpreter can then choose
99 what higher order action to take based on this value. For example, the simplest use
100 of the `hint` opcode is the ordinary “debug printf”. A more complex use of the opcode
101 allows Urbit to pass hints to Vere which produces effects on Earth (send a network
102 packet, get a keyboard input, etc.); bridging the Earth/Mars gap and allowing for vanes
103 to interact outside of Mars.

104 Key to Vere’s functionality is how it processes events. Events from the Unix en-
105 vironment are passed to Vere, which then translates and injects them into Arvo. This
106 mechanism demonstrates a profound integration between Vere and Arvo, where Vere
107 possesses direct knowledge of Arvo’s state. Notably, Vere can “scry” or query Arvo’s
108 state without altering it, retrieving information as needed. This capability allows Vere
109 to maintain the integrity of Arvo’s deterministic model while enabling dynamic in-

110 teraction with the outside world.

111 Urbit’s use of Unix as its BIOS is an important aspect of its runtime operation. By
112 leveraging the robust, widely-used Unix system, Vere ensures that Urbit can run on
113 a broad range of hardware platforms, effectively making Urbit’s Martian technology
114 accessible and operable in Earth’s diverse computing environments. This strategic use
115 of Unix not only provides the necessary I/O and optimizations for Arvo but also en-
116 sures that Urbit’s advanced and unique software architecture can function in tandem
117 with the existing, established hardware and software ecosystems of Earth.

118 Moreover, Vere’s design allows it to inject events into Arvo, a feature crucial for
119 maintaining the fluid communication and interaction between the Urbit system and its
120 underlying hardware and software infrastructure. This bidirectional communication
121 channel ensures that Arvo can respond to external stimuli while remaining true to its
122 functional and deterministic nature.

123 Thus, Vere serves as more than just a bridge between Martian technology and
124 Earthly hardware; it is a finely tuned conduit, optimized to uphold and facilitate the
125 unique computational paradigm that Arvo embodies.

126 3 %khan - The High-Level Thread Interface of Urbit

127 %khan, as a crucial component of Urbit’s architecture, functions as a high-level thread
128 interface, pivotal for managing both internal and external communications within the
129 Urbit ecosystem. It is designed to efficiently handle complex I/O operations, employ-
130 ing an I/O monad coupled with an exception monad. This approach allows %khan to
131 adeptly manage complex I/O tasks while effectively handling potential failures.

132 At its core, %khan’s design revolves around the use of ‘threads’, which are monadic
133 functions designed to handle arguments and produce results. These threads are opti-
134 mized for I/O operations, leveraging the I/O monad for structured and efficient task
135 handling. The exception monad further ensures robust management of potential fail-
136 ures and unexpected events. The weakness of threads is that it’s impermanent and
137 may fail unexpectedly. Since it is usually relying on some external In most of its inter-
138 mediate states, it expects only a small number of events (usually one), so if it receives
139 anything it didn’t expect, it fails. When code is upgraded, it’s impossible to upgrade
140 a running thread, so it fails.

141 %khan was conceived as a way to control Urbit ships from the exterior using
142 threads. The concept evolved a fair bit from proposal to implementation. In practice,
143 %khan is essentially an interface wrapper for Spider-based threads, which produces a
144 somewhat strange (but not unprecedented) situation in which a vane relies on a piece
145 of userspace infrastructure to function correctly. %spider is an agent for transient
146 thread-level operations.

147 %khan allows for pre-written threads that allow for easy hosting and maintenance
148 to be bundled and distributed; helping hosting companies and self-hosted ships to
149 easily run efficiently.

4 `conn.c`

`conn.c`¹ is a driver in Vere. It is a part of the “King” process. It exposes a Unix domain socket² at `/path/to/pier/.urb/conn.sock` for sending/receiving data from external processes.

The functionality of `conn.c`, particularly its ability to dispatch messages and handle various types of requests, is critical for the command of Urbit’s control plane. It serves as a way for Unix processes to receive insights about what is happening on Urbit. It also allows for pass commands either `%khan` and the injection of raw kernel moves (events) into the system.

From a technical perspective, `conn.c` accepts newt-encoded `++jamed` nouns (defined below) which take the form `[request-id command arguments]`. The newt-encoded format is:

`V.BBBB.JJJJ.JJJJ...`

Where

- `V` is the version (currently 0)
- `B` is the size of the `++jamed` noun in bytes (little endian)
- `J` is the `++jamed` noun (little endian)

This structure allows for a variety of commands to be executed, including:

1. `%ovum` - the injection of raw kernel moves
2. `%fyrd` - a direct shortcut to `%khan` commands
3. `%urth` - runtime subcommands like `%pack` or `%meld`
4. `%peek` - namespace scry requests into Arvo
5. `%peel` - emulated namespace scry requests into Vere

A valid `conn.c` command produces a newt-encoded `++jamed` noun with type `[request-id output]`, where:

- `request-id` matches the input `request-id`
- `output` depends on the command

An invalid `conn.c` command produces a newt-encoded jammed noun with type `[0 %bail error-code error-string]`.

¹<https://github.com/urbit/vere/blob/develop/pkg/vere/io/conn.c>

²https://en.wikipedia.org/wiki/Unix_domain_socket

5 %lick - The Low Level IPC Interface

Although also dealing with interprocess communication, %lick (UIP-101) was designed for a very different use-case than %khan: to allow external processes, in particular hardware drivers, to intercommunicate with Urbit. %lick focuses on creating a generic noun interface over domain sockets.

%lick manages IPC ports, and the communication between Urbit applications and POSIX applications via these ports. Other vanes and applications ask %lick to open an IPC port, notify it when something is connected or disconnected, and transfer data between itself and the Unix application. %lick works by opening a Unix socket for a particular process, which allows serialized IPC communications. These involve a jammed noun so the receiving process needs to know how to communicate in nouns. The IPC ports %lick creates are Unix domain sockets (AF_UNIX address family) of the SOCK_STREAM type. The connections are made via file path in .urb/dev of the pier.

The process on the host OS must therefore strip the first 5 bytes, ++cue³ the ++jammed noun, check the mark and (most likely) convert the noun into a native data structure.

To understand what %lick is doing, we need to look at Unix's IPC model briefly. IPC ("interprocess communication") describes any way that two processes in an operating system's shared context have to communicate with each other. %lick focuses on Unix domain sockets⁴, which are just communication endpoints⁵.

For instance, a valid use of %lick would use cards that look like this:

```
201 ++ init [[%pass / %arvo %l %spin /control]~ this]
202 ::
203 ++ on-arvo
204   |= [=wire =sign-arvo]
205   ?+ sign-arvo (on-arvo: def wire sign-arvo)
206     [%lick %soak *]
207     ?+ mark.sign-arvo [~ this]
208     ::
209       %connect
210       ~& > "connect"
211       :- this
212         [%pass /spit %arvo %l %spit /control %init area.state]~
213       ==
214     ==
215   ::
216 ++ send-state
217   |= =state
218   ^- card:agent:gall
219   :* %pass
```

³<https://docs.urbit.org/language/hoon/reference/stdlib/2p#cue>

⁴https://en.wikipedia.org/wiki/Unix_domain_socket

⁵<https://man7.org/linux/man-pages/man7/unix.7.html>

```

220     /spit
221     %arvo %l
222     %spit
223     /control
224     %state
225     [slick:state face.state food.state live.state]
226     ==

```

227 The vane definition of %lick is even simpler than %khan: it has no ++abet core
 228 and primarily communicates to the unix-duct in its state. The +\$owner is a +\$duct to
 229 handle the return %soak.

230 Arvo will send three types of %soaks to an open %lick port. These %soaks are
 231 %connect when the first connection is established on the Earth-side IPC port, %dis-
 232 connect when the last connection is broken on the Earth-side IPC port, and %error
 233 when an error occurs. %lick also will send a %disconnect %soak to every agent when
 234 Vere is started.

235 Gall needs to wrap %soak and %spit to route properly (see e.g. %++ap-generic-
 236 take). This lets multiple agents share sockets with the same name, while placing an
 237 agents sockets in a unique file path.

238 6 Conclusion

239 This exploration into Urbit's IPC interfaces, specifically %khan and %lick, has illumi-
 240 nated the profound architectural innovations that distinguish Urbit from traditional
 241 computing paradigms. By delving into the inner workings of Arvo and Vere, as well
 242 as the mechanisms facilitating communication between these components, we've un-
 243 covered a realm where simplicity and determinism reign supreme. This study not
 244 only demystifies the complex underpinnings of Urbit but also highlights its potential
 245 to redefine our interaction with digital ecosystems.

246 Urbit's design philosophy, rooted in the principles of a more principled and co-
 247 herent digital universe, challenges us to reconsider the foundations upon which our
 248 current software ecosystems are built. The contrast between the evolved, patchwork
 249 nature of Earth's software and the Maxwellian purity of Urbit's architecture serves as
 250 a poignant reminder of the potential for innovation in software design. As we stand
 251 at the crossroads of Earth's complexity and Mars' simplicity, the lessons drawn from
 252 Urbit's IPC interfaces beckon us towards a future where software is not only a tool
 253 but a reflection of a more streamlined and purposeful digital existence.

254 The journey through Urbit's landscape offers a vision of what digital computation
 255 could become: a domain where efficiency, clarity, and integrity are not aspirational
 256 but foundational. As we continue to navigate the evolving landscape of software de-
 257 velopment, the insights garnered from this analysis of Urbit will undoubtedly inspire
 258 further exploration and innovation. In embracing the principles elucidated by Urbit's
 259 design, we pave the way for a future where the digital and human experiences are
 260 more intimately aligned, fostering a computing environment that is both revolution-
 261 ary and deeply resonant with the principles of simplicity and purpose that guide our

262 exploration of the digital frontier.